

MUHUMMAD
HABIB

DATA SCIENCE PROJECT

MEDICAL INSURANCE COST PREDICTION

Kaggle

Github

@Linkedin:
Muhummad-habib

Medical Health Insurance Cost Prediction | Python + Regression Models

Medical Health Insurance Cost Prediction with python using Different Regression Models.

Overview

This project focuses on predicting medical health insurance costs using various regression models. The goal is to create a model that can accurately estimate the insurance costs for individuals based on their attributes. The project employs exploratory data analysis, data preprocessing, and several regression techniques to achieve this.

Table of Contents

- [Project Description](#)
- [Requirements](#)
- [Dataset](#)
- [Data Preprocessing](#)
- [Exploratory Data Analysis](#)
- [Regression Models](#)
- [Model Evaluation](#)
- [Key Insights](#)
- [Suggestions](#)
- [Conclusion](#)

Project Description

In this project, we aim to predict medical health insurance costs for individuals based on various factors such as age, sex, BMI, number of children, smoking habits, and region. The project involves the following steps:

1. **Data Loading and Overview:** The project begins by loading the dataset containing information about individuals and their insurance costs.
2. **Exploratory Data Analysis:** We analyze the dataset to gain insights into the data distribution, relationships between variables, and identify potential patterns.
3. **Data Preprocessing:** Data preprocessing steps are performed, including handling missing values, encoding categorical variables, and scaling numerical features.
4. **Regression Models:** Several regression models are implemented, including Multiple Linear Regression, Lasso Regression, Ridge Regression, ElasticNet Regressor, Random Forest Regressor, and Polynomial Regression.
5. **Model Evaluation:** The models are evaluated using metrics such as Mean Absolute Error, Mean Squared Error, and Root Mean Squared Error. The performance of each model is compared.
6. **Conclusion:** The project concludes with a summary of key insights obtained from the analysis and suggestions for further improvement.

Requirements

1. **Python**
2. **Integrated Development Environment (IDE):** You can use IDEs like **Jupyter Notebook**, **PyCharm**, **Visual Studio Code**, or any other of your preference.
3. **Dataset:** Collect the dataset containing medical records, including factors like age, gender, BMI, number of dependents, smoking status, region, and charges.

4. **Data Preprocessing Libraries:** You might need libraries like **Pandas** for data manipulation and cleaning, and **NumPy** for numerical operations.
5. **Data Visualization Libraries:** Consider using libraries like **Matplotlib** and **Seaborn** for visualizing the data distribution, trends, and anomalies.
6. **Statistical Libraries:** Depending on your project's needs, **Scipy** might be useful for advanced statistical operations.
7. **Machine Learning Libraries:** You will likely need **Scikit-learn** for standard machine learning algorithms and **TensorFlow** or **PyTorch** for any deep learning implementations.
8. **Regression Algorithms:** Implementing a project on insurance cost prediction will require a variety of regression algorithms such as Linear Regression, **Ridge** and **Lasso** Regression, **Decision Tree** Regression, **Random Forest** Regression, and possibly more advanced techniques based on your research.
9. **Model Evaluation Tools:** Libraries like **Scikit-learn** contain useful functions for evaluating your regression models, including metrics like Mean Absolute Error (MAE), Mean Squared Error (MSE), and **R-squared**.
10. **Data Preprocessing Tools:** You might need tools like **Scikit-learn's** Pipeline and preprocessing module for efficiently handling data preprocessing tasks.
11. **Model Deployment (Optional):** If you plan to deploy the trained models, you might need libraries like **Flask** or **FastAPI** for creating APIs.

Dataset

The dataset used for this project contains information about individuals' attributes and their corresponding medical insurance costs. The attributes include age, sex, BMI, number of children, smoking habits, and region. The dataset is loaded and explored using Python's pandas library.

Dataset Source: <https://www.kaggle.com/datasets/teertha/ushealthinsurancedataset>

```
# !pip install sklearn
```

```
# !pip install scikit-learn
```

```
import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
import plotly.express as px
import plotly.graph_objects as go
import plotly.figure_factory as ff
from plotly.offline import download_plotlyjs, init_notebook_mode, iplot
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
# !pip install plotly
# !pip install plotly
```

```
df = pd.read_csv('insurance.csv')
```

Data Preprocessing

Data preprocessing involves handling missing values, encoding categorical variables, and scaling numerical features. Label encoding and one-hot encoding are used for categorical variables. Additionally, feature scaling is performed to ensure that numerical features are on similar scales.

```
df.head()
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
df.describe
```

```
<bound method NDFrame.describe of
region      charges
0      19  female  27.900      0  yes  southwest  16884.92400
1      18   male  33.770      1   no  southeast   1725.55230
2      28   male  33.000      3   no  southeast   4449.46200
3      33   male  22.705      0   no  northwest  21984.47061
4      32   male  28.880      0   no  northwest   3866.85520
...
1333   50   male  30.970      3   no  northwest  10600.54830
1334   18  female  31.920      0   no  northeast   2205.98080
1335   18  female  36.850      0   no  southeast   1629.83350
1336   21  female  25.800      0   no  southwest   2007.94500
1337   61  female  29.070      0  yes  northwest  29141.36030
```

```
[1338 rows x 7 columns]>
```

```
df.describe()
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
df.isnull().sum()
```

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

there's no any null values


Exploratory Data Analysis

Exploratory Data Analysis (EDA) is performed to understand the distribution of variables and identify relationships. Visualizations using libraries such as seaborn, matplotlib, and plotly are used to create various plots and graphs.

#WE CAN APPLY LOG TRANSFORM TO CORREST SKEWNESS


```
skewed = sns.displot(df['charges']) #Charges is right-skewed
skewed
```

```
<seaborn.axisgrid.FacetGrid at 0x1f0a4f06040>
```



```
log_trans = sns.displot(np.log10(df['charges'])) #skewness is corrected using log
log_trans
```

```
<seaborn.axisgrid.FacetGrid at 0x1f0ad223f10>
```



#CHARGES BY REGION

```
charges = df['charges'].groupby(df['region']).sum().sort_values(ascending = True)
charges
```

```
region
southwest      4.012755e+06
northwest      4.035712e+06
northeast      4.343669e+06
southeast      5.363690e+06
Name: charges, dtype: float64
```

Plotting charges by region with plotly.express aka px

```
fig = px.bar(charges, title='Charges by Region', color=charges,
color_continuous_scale='plasma')
fig.update_layout(
    margin=dict(t=50, b=0, l=0, r=0),
    titlefont=dict(size=20),
    xaxis_tickangle=-45, # Specify the angle at which x-axis labels are
displayed
)
fig.update_yaxes(showticklabels=False, title='')
fig.update_xaxes(title='')
fig.update_traces(
    texttemplate='%{y}', # Display y values as text on the bars
    textposition='outside', # Position the text outside the bars
    hovertemplate='<b>{x}</b><br>Charges: %{y}', # Customize hover template
)
fig.show()
```



```

{"config":{"plotlyServerURL":"https://plot.ly"},"data":
[{"alignmentgroup":"True","hovertemplate":"<b>{%x}</b><br>Charges: %
{y}","legendgroup":"","marker":{"color":
[4012754.64762,4035711.99654,4343668.583309,5363689.76329],"coloraxis":"coloraxis
","pattern":
{"shape":"","name":"","offsetgroup":"","orientation":"v","showlegend":false,"te
xtposition":"outside","texttemplate":"%{y}","type":"bar","x":
["southwest","northwest","northeast","southeast"],"xaxis":"x","y":
[4012754.64762,4035711.99654,4343668.583309,5363689.76329],"yaxis":"y"}],"layout"
:{"barmode":"relative","coloraxis":{"colorbar":{"title":
{"text":"color"}}, "colorscale":[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],[1,"#f0f921"]]},"legend":
{"tracegroupgap":0,"margin":{"b":0,"l":0,"r":0,"t":50},"template":{"data":
{"bar":[{"error_x":{"color":"#2a3f5f"},"error_y":{"color":"#2a3f5f"},"marker":
{"line":{"color":"#E5ECF6"},"width":0.5},"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2}}, "type":"bar"}],"barpolar":
[{"marker":{"line":{"color":"#E5ECF6"},"width":0.5},"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2}}, "type":"barpolar"}],"carpet":
[{"aaxis":
{"endlinecolor":"#2a3f5f","gridcolor":"white","linecolor":"white","minorgridcolor
":"white","startlinecolor":"#2a3f5f"},"baxis":
{"endlinecolor":"#2a3f5f","gridcolor":"white","linecolor":"white","minorgridcolor
":"white","startlinecolor":"#2a3f5f"},"type":"carpet"}],"choropleth":
[{"colorbar":{"outlinewidth":0,"ticks":""},"type":"choropleth"}],"contour":
[{"colorbar":{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"contour"}],"contourcarpet":[{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"contourcarpet"}],"heatmap":[{"colorbar":
{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"heatmap"}],"heatmapgl":[{"colorbar":
{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"heatmapgl"}],"histogram":[{"marker":{"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2}}, "type":"histogram"}],"histogram2
d":[{"colorbar":{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"histogram2d"}],"histogram2dcontour":[{"colorbar":
{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"histogram2dcontour"}],"mesh3d":[{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"mesh3d"}],"parcoords":[{"line":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"parcoords"}],"pie":

```

```
[{"automargin":true,"type":"pie"},"scatter":[{"fillpattern":
{"fillmode":"overlay","size":10,"solidity":0.2},"type":"scatter"}],"scatter3d":
[{"line":{"colorbar":{"outlinewidth":0,"ticks":""},"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scatter3d"},"scattercarpet":[{"marker":
{"colorbar":{"outlinewidth":0,"ticks":""},"type":"scattercarpet"},"scattergeo":
{"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scattergeo"},"scattergl":[{"marker":
{"colorbar":{"outlinewidth":0,"ticks":""},"type":"scattergl"},"scattermapbox":
{"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scattermapbox"},"scatterpolar":
{"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scatterpolar"},"scatterpolargl":
{"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scatterpolargl"},"scatterternary":
{"marker":{"colorbar":
{"outlinewidth":0,"ticks":""},"type":"scatterternary"},"surface":[{"colorbar":
{"outlinewidth":0,"ticks":""},"colorscale":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"surface"},"table":[{"cells":{"fill":
{"color":"#EBF0F8"},"line":{"color":"white"},"header":{"fill":
{"color":"#C8D4E3"},"line":{"color":"white"},"type":"table"}]}],"layout":
{"annotationdefaults":
{"arrowcolor":"#2a3f5f","arrowhead":0,"arrowwidth":1,"autotypenumbers":"strict",
"coloraxis":{"colorbar":{"outlinewidth":0,"ticks":""},"colorscale":{"diverging":
[[0,"#8e0152"],[0.1,"#c51b7d"],[0.2,"#de77ae"],[0.3,"#f1b6da"],[0.4,"#fde0ef"],
[0.5,"#f7f7f7"],[0.6,"#e6f5d0"],[0.7,"#b8e186"],[0.8,"#7fbcb41"],[0.9,"#4d9221"],
[1,"#276419"]],"sequential":[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],[1,"#f0f921"]],"sequentialminus":[[0,"#0d0887"],
[0.1111111111111111,"#46039f"],[0.2222222222222222,"#7201a8"],
[0.3333333333333333,"#9c179e"],[0.4444444444444444,"#bd3786"],
[0.5555555555555556,"#d8576b"],[0.6666666666666666,"#ed7953"],
[0.7777777777777778,"#fb9f3a"],[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]]}],"colorway":
["#636efa","#EF553B","#00cc96","#ab63fa","#FFA15A","#19d3f3","#FF6692","#B6E880",
"#FF97FF","#FECB52"],"font":{"color":"#2a3f5f"},"geo":
{"bgcolor":"white","lakecolor":"white","landcolor":"#E5ECF6","showlakes":true,"sh
owland":true,"subunitcolor":"white"},"hoverlabel":
{"align":"left"},"hovermode":"closest","mapbox":
{"style":"light"},"paper_bgcolor":"white","plot_bgcolor":"#E5ECF6","polar":
{"angularaxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"bgcolor":"#E5ECF6","radiala
xis":{"gridcolor":"white","linecolor":"white","ticks":""},"scene":{"xaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"linecolor":"white",
"showbackground":true,"ticks":"","zerolinecolor":"white"},"yaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"linecolor":"white",
"showbackground":true,"ticks":"","zerolinecolor":"white"},"zaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"linecolor":"white",
"showbackground":true,"ticks":"","zerolinecolor":"white"}},"shapedefaults":
{"line":{"color":"#2a3f5f"},"ternary":{"aaxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"baxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"bgcolor":"#E5ECF6","caxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"title":{"x":5.0e-
2},"xaxis":
{"automargin":true,"gridcolor":"white","linecolor":"white","ticks":"","title":
{"standoff":15},"zerolinecolor":"white","zerolinewidth":2},"yaxis":
```




```
{ "automargin": true, "gridcolor": "white", "linecolor": "white", "ticks": "", "title": { "standoff": 15 }, "zerolinecolor": "white", "zerolinewidth": 2 } }, "title": { "font": { "size": 20 }, "text": "Charges by Region", "xaxis": { "anchor": "y", "domain": [ 0, 1 ], "tickangle": -45, "title": { "text": " " } }, "yaxis": { "anchor": "x", "domain": [ 0, 1 ], "showticklabels": false, "title": { "text": " " } } }
```

#WE CHECK THE CHARGES BY REGION WHO ARE SMOKERS


```
sns.barplot(data = df, x = 'region', y = 'charges', hue = 'smoker', palette = 'rocket' )
```

<Axes: xlabel='region', ylabel='charges'>



```
plt.figure(figsize=(12,6))
sns.violinplot(data = df, x = 'children', y = 'charges', hue = 'smoker', split = False, palette = 'magma')
```


<Axes: xlabel='children', ylabel='charges'>



#WE CHECK THE CHARGES BY REGION BY THIER GENDER

```
sns.barplot(data = df, x='region', y='charges', hue='sex', palette='spring')
```


<Axes: xlabel='region', ylabel='charges'>



#WE CHECK THE CHARGES BY REGION BY THE CHILDREN AVAILABLE

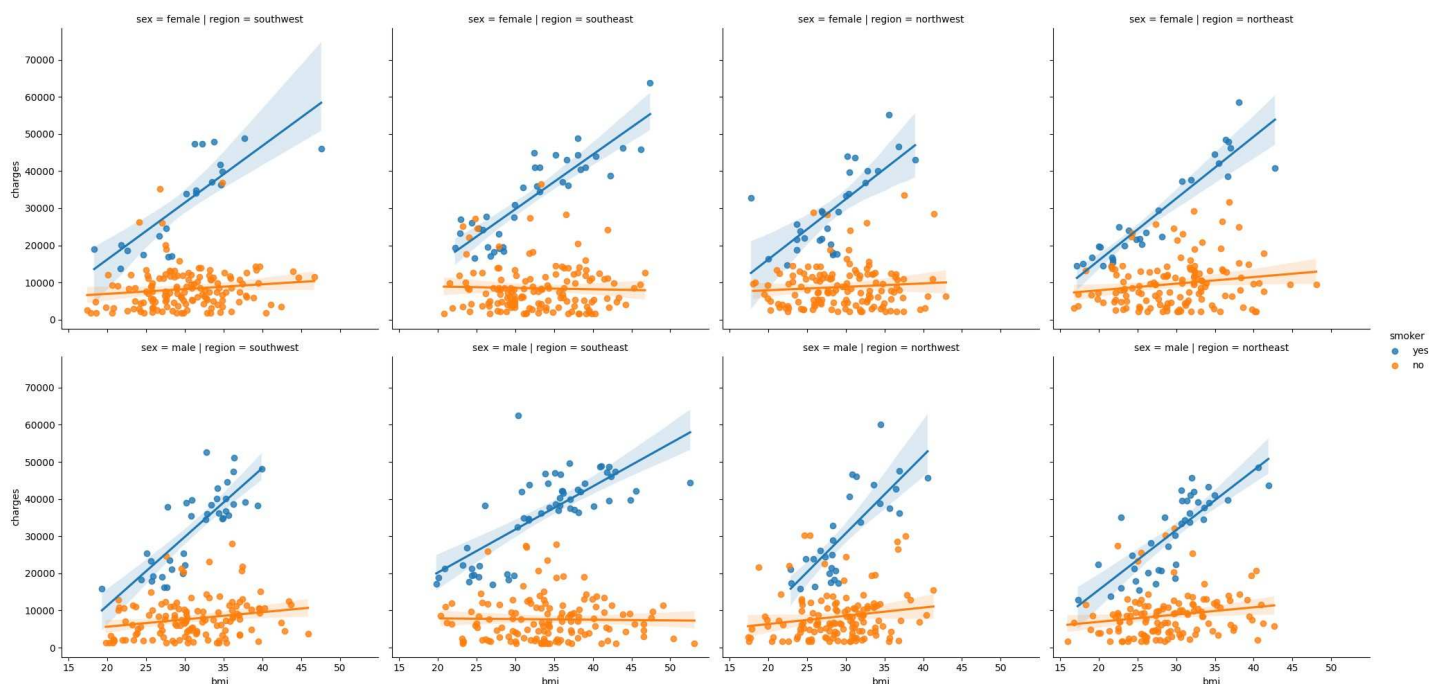
```
plt.figure(figsize=(12,6))
sns.barplot(data = df, x = 'region', y = 'charges', hue = 'children',
palette='Paired')
```

<Axes: xlabel='region', ylabel='charges'>



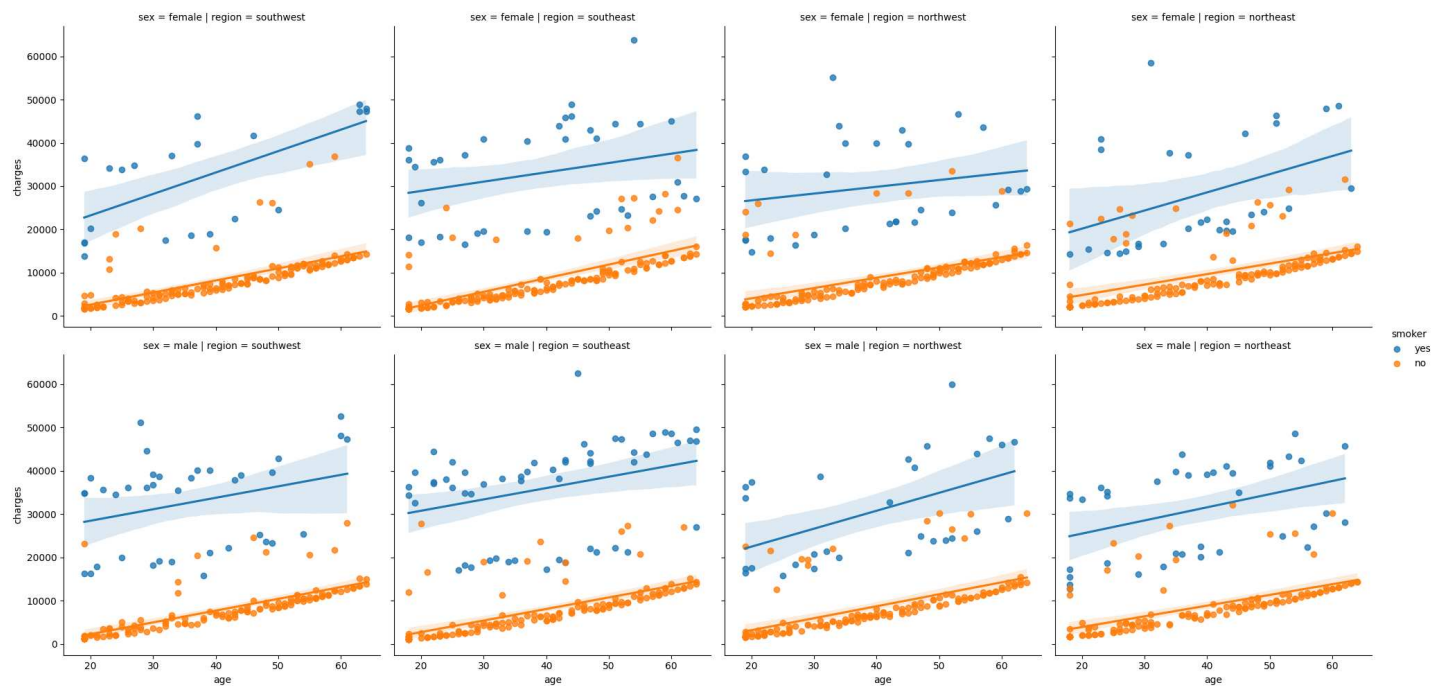
```
#regression plot to understand the relationship between the bmi and charges
considering
sns.lmplot(x = "bmi", y = "charges", row = "sex", col = "region", hue = 'smoker',
data = df)
```

```
<seaborn.axisgrid.FacetGrid at 0x1f0ba646e50>
```



```
#regression plot to understand the relationship between the Age and Charges
considering
sns.lmplot(x = "age", y = "charges", row = "sex", col = "region", hue = 'smoker',
data = df)
```

```
<seaborn.axisgrid.FacetGrid at 0x1f0bc80bf40>
```



From the chart above we can see that those who are smokers get a higher medical insurance charges than those that are none smoker

```
# Convert object labels to categorical data type
df[['sex', 'region', 'smoker']] = df[['sex', 'region',
'smoker']].astype('category')

# Check the updated data types
print(df.dtypes)

age          int64
sex          category
bmi         float64
children     int64
smoker       category
region       category
charges     float64
dtype: object

# Converting category labels into numerical using LabelEncoder
from sklearn import preprocessing
label = preprocessing.LabelEncoder()

label.fit(df.sex.drop_duplicates())
df.sex = label.transform(df.sex)

label.fit(df.smoker.drop_duplicates())
df.smoker = label.transform(df.smoker)

label.fit(df.region.drop_duplicates())
df.region = label.transform(df.region)

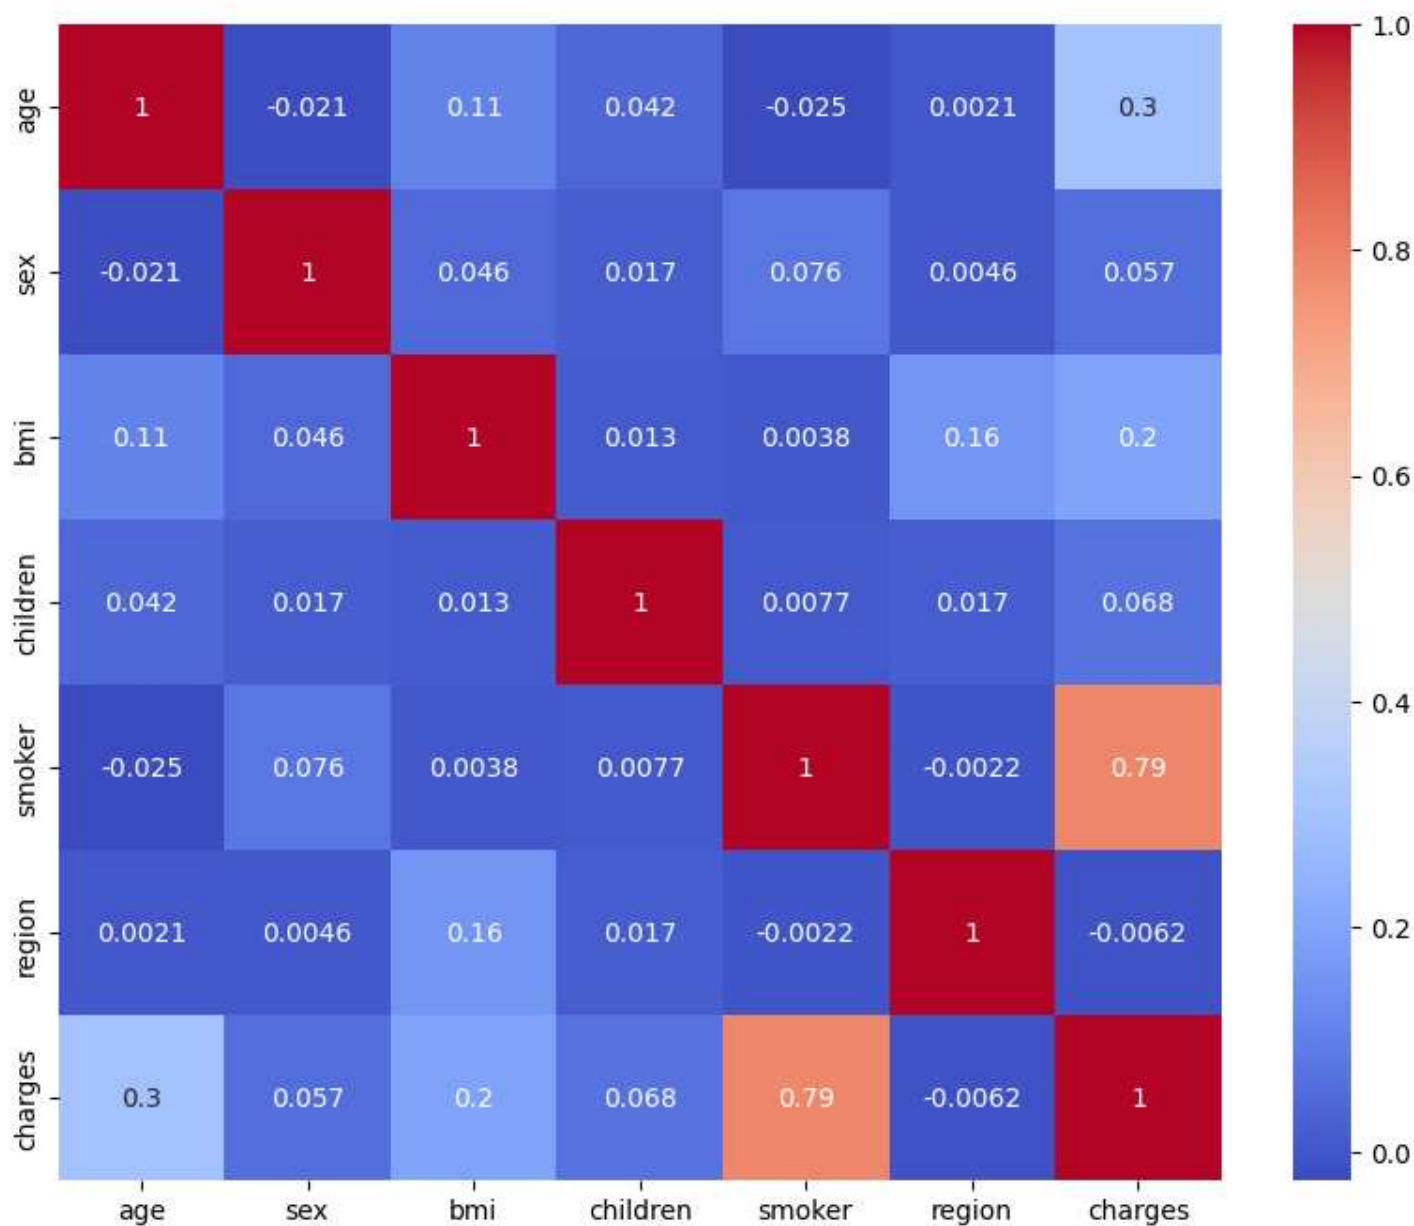
df.head()
```

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200

```
3  33    1  22.705      0      0      1  21984.47061
4  32    1  28.880      0      0      1   3866.85520
```

```
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(),cmap='coolwarm',annot=True)
```

<Axes: >



#we split our model

```
x = df.drop(['charges'], axis = 1)
y = df['charges']
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x, y, test_size = 0.2,
random_state = 42)
```

!pip install statsmodels

```
import statsmodels.api as sm #WE GET THE STATISTICAL MODEL
#add constant to predictor variables
x2 = sm.add_constant(x_train)
```



```
#fit linear regression model
model = sm.OLS(y_train, x2).fit()
```

```
model.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

Dep. Variable:	charges	R-squared:	0.742			
Model:	OLS	Adj. R-squared:	0.740			
Method:	Least Squares	F-statistic:	508.7			
Date:	Sun, 20 Aug 2023	Prob (F-statistic):	2.71e-308			
Time:	21:17:48	Log-Likelihood:	-10845.			
No. Observations:	1070	AIC:	2.170e+04			
Df Residuals:	1063	BIC:	2.174e+04			
Df Model:	6					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	-1.195e+04	1086.938	-10.991	0.000	-1.41e+04	-9813.820
age	257.0563	13.452	19.109	0.000	230.661	283.451
sex	-18.7915	375.770	-0.050	0.960	-756.126	718.543
bmi	335.7815	31.655	10.607	0.000	273.668	397.895
children	425.0915	154.431	2.753	0.006	122.067	728.116
smoker	2.365e+04	465.245	50.829	0.000	2.27e+04	2.46e+04
region	-271.2843	170.373	-1.592	0.112	-605.590	63.022
=====				=====		
Omnibus:	251.823	Durbin-Watson:	2.085			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	611.548			
Skew:	1.251	Prob(JB):	1.60e-133			
Kurtosis:	5.731	Cond. No.	299.			
=====				=====		

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

Standard Errors:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

R-squared: 0.747

- The coefficient of determination.
- Proportion of variance in the response variable explained by predictor variables.
- 74.7% of the variation in exam scores explained by hours studied and prep exams taken.

F-statistic: 523.6

- Overall F-statistic for the regression model.

Prob (F-statistic): 3.16e-313

- p-value associated with the overall F-statistic.
- Determines if the regression model is statistically significant.
- p-value < 0.05 indicates a significant association of predictor variables with the response variable.

coef:

- Coefficients for each predictor variable.
- Represents average expected change in the response variable, with other predictors constant.

P>|t|:

- Individual p-values for each predictor variable.
- Indicates statistical significance.
- Statistically significant at $\alpha < 0.05$, not significant at $\alpha \geq 0.05$.
- Example: "sex" not statistically significant, may be removed from the model.

Regression Models

Multiple regression models are implemented to predict insurance costs.

MULTIPLE LINEAR REGRESSION MODEL

```
from sklearn.linear_model import LinearRegression
```

```
lm = LinearRegression()
```

```
lm.fit(x_train,y_train)
```

```
LinearRegression()
```

```
print('Intercept', lm.intercept_)
print('Coefficient', lm.coef_)
print('Score', lm.score(x_test, y_test))
```

```
Intercept -11946.60656726302
```

```
Coefficient [ 2.57056264e+02 -1.87914567e+01  3.35781491e+02  4.25091456e+02
 2.36478181e+04 -2.71284266e+02]
```

```
Score 0.7833463107364539
```

```
coeff_df = pd.DataFrame(lm.coef_,x.columns,columns=['Coefficient'])
coeff_df
```

	Coefficient
age	257.056264
sex	-18.791457
bmi	335.781491
children	425.091456
smoker	23647.818096
region	-271.284266

#WE PREDICT OUR MODEL

```
y_pred1 = lm.predict(x_test)
```

```
plt.figure(figsize=(10, 5))
```

```
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")
```

```
sns.distplot(y_pred1, hist=False, color="b", label="Fitted Values" , ax=ax1)
```


```
plt.title('Actual vs Fitted Values for Charges')
```

```
plt.xlabel('Charges')
```

```
plt.ylabel('')
```


```
plt.show()
```

```
plt.close()
```



ANOTHER PLOTTING TECHNIQUE

```
plt.figure(figsize=(10, 5))
# actual values
plt.plot([i for i in range(len(y_test))], np.array(y_test), c='g', label="actual values")
# predicted values
plt.plot([i for i in range(len(y_test))], y_pred1, c='m', label="predicted values")
plt.legend()
plt.show()
```



#DATAFRAME FOR ACTUAL AND PREDICTED VALUE

```
predicted1 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred1})
predicted1.head()
```

	Actual	Predicted
764	9095.06825	8924.407244
887	5272.17580	7116.295018
890	29330.98315	36909.013521
1293	9301.89355	9507.874691
259	33750.29180	27013.350008

LASSO REGRESSION MODEL

```
from sklearn.linear_model import LassoCV
from sklearn.model_selection import RepeatedKFold

#define cross-validation method to evaluate model
cv = RepeatedKFold(n_splits=10, n_repeats=4, random_state=101) #we'll use the RepeatedKF
#define model
lasso_model = LassoCV(alphas=(0.1, 1.0, 10.0), cv=cv, n_jobs=-1)
#fit model
lasso_model.fit(x_train, y_train)

LassoCV(alphas=(0.1, 1.0, 10.0),
        cv=RepeatedKFold(n_repeats=4, n_splits=10, random_state=101),
        n_jobs=-1)

print('Intercept', lm.intercept_)
print('Coefficient', lm.coef_)
print('Score', lm.score(x_test, y_test))

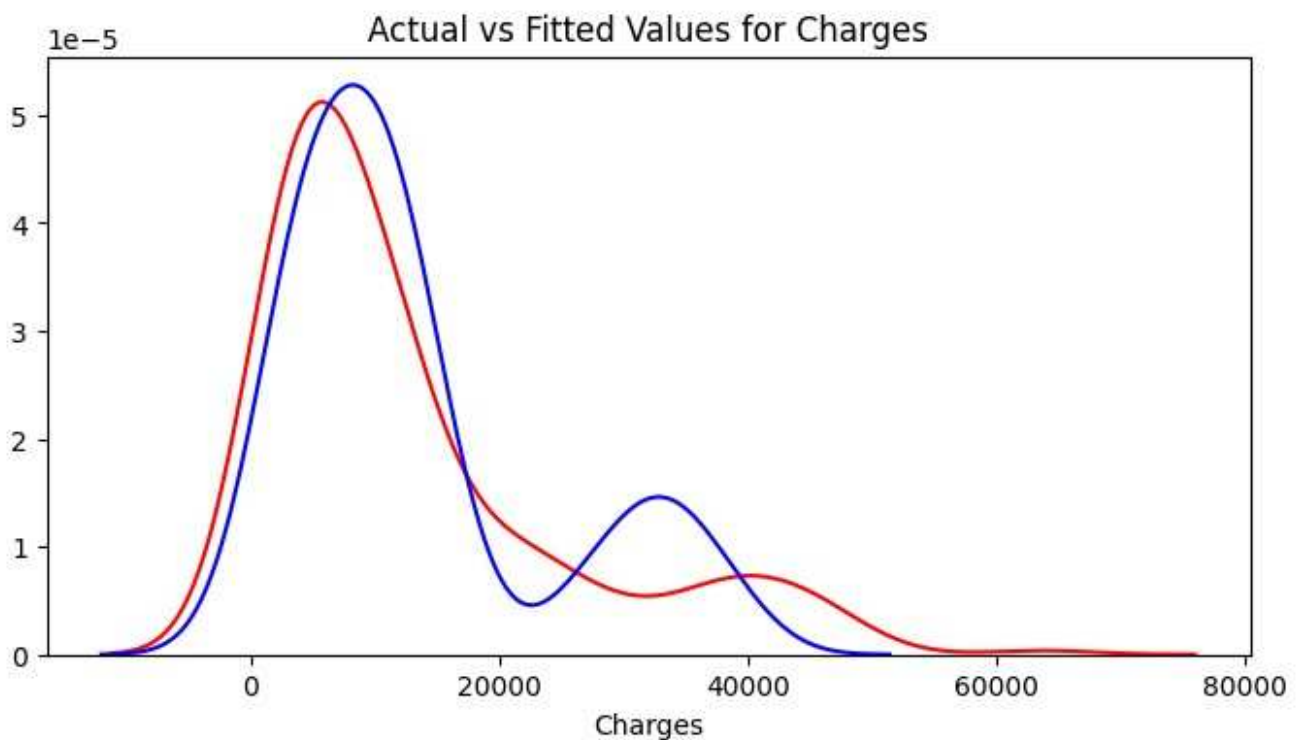
Intercept -11946.60656726302
Coefficient [ 2.57056264e+02 -1.87914567e+01  3.35781491e+02  4.25091456e+02
 2.36478181e+04 -2.71284266e+02]
Score 0.7833463107364539
```

#WE PREDICT OUR LASSOCV REGRESSION MODEL

```
y_pred2 = lasso_model.predict(x_test)
```

Visualising the Lasso Regression results

```
plt.figure(figsize=(8, 4))
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")
sns.distplot(y_pred2, hist=False, color="b", label="Fitted Values" , ax=ax1)
plt.title('Actual vs Fitted Values for Charges')
plt.xlabel('Charges')
plt.ylabel('')
plt.show()
plt.close()
```



#DATAFRAME FOR ACTUAL AND PREDICTED VALUE

```
predicted2 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred2})
predicted2.head()
```

	Actual	Predicted
764	9095.06825	8911.659191
887	5272.17580	7123.036794
890	29330.98315	36852.479165
1293	9301.89355	9515.428284
259	33750.29180	26976.771303

RIDGE REGRESSION MODEL

```
from sklearn.linear_model import Ridge
from sklearn.linear_model import RidgeCV
from sklearn.model_selection import RepeatedKFold
```

```
#define cross-validation method to evaluate model
# cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=101)
```

```
#define model
rid_model = RidgeCV(alphas=(0.1, 1.0, 10.0), cv=cv,
scoring='neg_mean_absolute_error')
#fit model
rid_model.fit(x_train, y_train)
```

```
RidgeCV(cv=RepeatedKFold(n_repeats=4, n_splits=10, random_state=101),
scoring='neg_mean_absolute_error')
```

```
print(rid_model.intercept_)
print(rid_model.coef_)
print(rid_model.score(x_test, y_test))
```

```
-11943.455858669686
[ 2.57035723e+02 -1.80109296e+01  3.35782790e+02  4.25136864e+02
 2.36341868e+04 -2.71303705e+02]
0.7833217621706015
```



```
#WE PREDICT OUR RIDGECV REGRESSION MODEL
```

```
y_pred3 = rid_model.predict(x_test)
```

```
# Visualising the Ridge Regression results
```

```
plt.figure(figsize=(10, 5))
```

```
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")
```

```
sns.distplot(y_pred3, hist=False, color="b", label="Fitted Values" , ax=ax1)
```


```
plt.title('Actual vs Fitted Values for Charges')
```

```
plt.xlabel('Charges')
```

```
plt.ylabel('')
```

```
plt.show()
```

```
plt.close()
```



```
#DATAFRAME FOR ACTUAL AND PREDICTED VALUE
```

```
predicted3 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred3})
```

```
predicted3.head()
```

	Actual	Predicted
764	9095.06825	8926.757144
887	5272.17580	7118.725826
890	29330.98315	36897.233767
1293	9301.89355	9511.011287
259	33750.29180	27003.281662

ELASTICNET REGRESSOR

```
from sklearn.model_selection import cross_val_score
```

```
from sklearn.model_selection import RepeatedKFold
```

```
from sklearn.linear_model import ElasticNet
```

```
#define cross-validation method to evaluate model
```

```
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
```

```
# define model
```

```
net_model = ElasticNet(alpha=0.1, l1_ratio=0.9, fit_intercept=True,  
max_iter=1000, random_state=1)
```

```
# evaluate model
scores = cross_val_score(net_model, x_train,y_train,
scoring='neg_mean_absolute_error', cv=cv)

net_model.fit(x_train, y_train)


ElasticNet(alpha=0.1, l1_ratio=0.9, random_state=1)

print(net_model.intercept_)
print(net_model.coef_)
print(net_model.score(x_test, y_test))

-11626.875213341189
[ 254.98398073   56.64093616  335.90294961  429.44085068
 22273.01597223 -273.05052278]
0.7789260673175409

y_pred4 = net_model.predict(x_test)
```

```
# Visualising the ElasticNet Regressor results
plt.figure(figsize=(9, 4))
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")
sns.distplot(y_pred4, hist=False, color="b", label="Fitted Values" , ax=ax1)
plt.title('Actual vs Fitted Values for Charges')
plt.xlabel('Charges')
plt.ylabel('')
plt.show()
plt.close()
```



```
#DATAFRAME FOR ACTUAL AND PREDICTED VALUE
predicted4 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred4})
predicted4.head()
```

	Actual	Predicted
764	9095.06825	9162.642377
887	5272.17580	7363.304117
890	29330.98315	35722.815803
1293	9301.89355	9822.122303
259	33750.29180	25996.448958

RANDOM FOREST REGRESSOR

```
from sklearn.ensemble import RandomForestRegressor
```

```
Ram_reg = RandomForestRegressor(n_estimators=100, criterion='squared_error',  
min_samples_split=2)  
# fit the regressor model  
Ram_reg.fit(x_train, y_train)
```

```
RandomForestRegressor()
```


```
print(Ram_reg.score(x_test, y_test))
```

```
0.8659967534430548
```

```
y_pred5 = Ram_reg.predict(x_test)
```

```
# Visualising the Random Forest Regressor results
```

```
plt.figure(figsize=(9, 4))  
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")  
sns.distplot(y_pred5, hist=False, color="b", label="Fitted Values" , ax=ax1)  
plt.title('Actual vs Fitted Values for Charges')  
plt.xlabel('Charges')  
plt.ylabel('')  
plt.show()  
plt.close()
```



```
#DATAFRAME FOR ACTUAL AND PREDICTED VALUE
```


```
predicted5 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred5})  
predicted5.head()
```

	Actual	Predicted
764	9095.06825	10348.293849
887	5272.17580	5268.427971
890	29330.98315	28489.521986
1293	9301.89355	9908.482924
259	33750.29180	34705.577697

WE TRY SELECT THE BEST FEATURES USING FEATURE IMPORTANCE FROM RANDOM FOREST REGRESSOR

```
features = x.columns
importances = Ram_reg.feature_importances_
indices = np.argsort(importances)
```

```
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='#8f63f4',
align='center')
plt.yticks(range(len(indices)), features[indices])
plt.xlabel('Relative Importance')
plt.show()
```



We can see that the smoker, bmi and age are more important features compared to the other features.

#We select the import features

```
x = df.drop(['charges', 'region', 'sex'], axis = 1)
y = df['charges']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
random_state = 0)
```

x

	age	bmi	children	smoker
0	19	27.900	0	1
1	18	33.770	1	0
2	28	33.000	3	0
3	33	22.705	0	0
4	32	28.880	0	0
...
1333	50	30.970	3	0
1334	18	31.920	0	0

1335	18	36.850	0	0
1336	21	25.800	0	0
1337	61	29.070	0	1

[1338 rows x 4 columns]

WE BUILD A MODEL USING THE POLYNOMIAL REGRESSION AFTER FEATURE IMPORTANCE

POLYNOMIAL REGRESSION MODEL

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly_reg = PolynomialFeatures(degree=2)
```

```
x_poly = poly_reg.fit_transform(x)
```

```
x_train,x_test,y_train,y_test = train_test_split(x_poly, y, test_size = 0.2,
random_state = 0)
```

```
pol_reg = LinearRegression()
pol_reg.fit(x_train, y_train)
```

```
LinearRegression()
```

```
print(pol_reg.intercept_)
print(pol_reg.coef_)
print(pol_reg.score(x_test, y_test))
```

```
-5325.881705251873
[ 0.00000000e+00 -4.01606591e+01  5.23702019e+02  8.52025026e+02
 -9.52698471e+03  3.04430186e+00  1.84508369e+00  6.01720286e+00
  4.20849790e+00 -9.38983382e+00  3.81612289e+00  1.40840670e+03
 -1.45982790e+02 -4.46151855e+02 -9.52698471e+03]
0.8812595703345225
```


#WE PREDICT OUR POLYNOMIAL REGRESSION MODEL

```
y_pred6 = pol_reg.predict(x_test)
```

Visualising the Polynomial Regression results

```
plt.figure(figsize=(9, 4))
ax1 = sns.distplot(y_test, hist=False, color="r", label="Actual Value")
sns.distplot(y_pred6, hist=False, color="b", label="Fitted Values" , ax=ax1)
plt.title('Actual vs Fitted Values for Charges')
plt.xlabel('Charges')
plt.ylabel('')
```

```
plt.show()
plt.close()
```

Model Evaluation

The performance of each regression model is evaluated using various metrics, including Mean Absolute Error (MAE), Mean Squared Error (MSE), and Root Mean Squared Error (RMSE). The evaluation results provide insights into the accuracy of the models in predicting insurance costs.

```
from sklearn import metrics
from sklearn.metrics import r2_score
```

MULTIPLE LINEAR REGRESSION

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred1))
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred1))
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,
y_pred1)))
```

Mean Absolute Error: 10976.93032415958
Mean Square Error: 225340349.96268988
Root Mean Square Error: 15011.340711698269

LASSOCV REGRESSION

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred2))
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred2))
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,
y_pred2)))
```

Mean Absolute Error: 10968.76909663481
Mean Square Error: 224915279.3845799
Root Mean Square Error: 14997.175713599541

RIDGECV REGRESSION

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred3))
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred3))
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,
y_pred3)))
```

Mean Absolute Error: 10975.462165520717
Mean Square Error: 225254498.69396165
Root Mean Square Error: 15008.480892280926

ELASTICNET REGRESSION

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred4))  
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred4))  
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,  
y_pred4)))
```

Mean Absolute Error: 10838.727060068726
Mean Square Error: 216982294.5125357
Root Mean Square Error: 14730.31888699412

RANDOM FOREST REGRESSOR

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred5))  
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred5))  
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,  
y_pred5)))
```

Mean Absolute Error: 11484.288418379512
Mean Square Error: 266813319.89767018
Root Mean Square Error: 16334.421321175421

POLYNOMIAL REGRESSION

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred6))  
print('Mean Square Error:', metrics.mean_squared_error(y_test, y_pred6))  
print('Root Mean Square Error:', np.sqrt(metrics.mean_squared_error(y_test,  
y_pred6)))
```

Mean Absolute Error: 2824.49504547766
Mean Square Error: 18895160.09878044
Root Mean Square Error: 4346.856346692451

From the model evaluation score, we can see that polynomial regression and the Random Forest Regressor are performing well than the other models.

Key Insights

1. **Feature Importance:** After analyzing the dataset, we found that the most influential factors affecting medical insurance charges are whether the person is a smoker, their BMI (Body Mass Index), and their age.
2. **Smoking Impact:** Smokers tend to have significantly higher medical insurance charges compared to non-smokers. This is a crucial insight that highlights the importance of lifestyle choices on healthcare costs.
3. **Age and Charges:** The age of an individual is positively correlated with insurance charges. Older individuals tend to have higher medical costs, which is understandable given the increased likelihood of health issues as age advances.
4. **BMI Influence:** Higher BMI values are associated with increased medical insurance charges. This suggests that maintaining a healthy weight can positively impact healthcare costs.
5. **Model Performance:** We evaluated multiple regression models including Multiple Linear Regression, LassoCV, RidgeCV, ElasticNet, Random Forest Regressor, and Polynomial Regression. Among these, Polynomial Regression and Random Forest Regressor showed better performance in predicting medical insurance charges.

6. **Polynomial Model Advantage:** The Polynomial Regression model showed the lowest Mean Absolute Error, Mean Square Error, and Root Mean Square Error among all models, indicating its capability to better capture the underlying patterns in the data.

Suggestions

1. **Promote Healthy Lifestyles:** Encourage smoking cessation programs and awareness campaigns to reduce the number of smokers. This could lead to substantial savings in healthcare costs.
2. **Health Awareness:** Run campaigns to educate people about the impact of BMI on health and insurance costs. Promoting healthy eating and exercise can help individuals maintain a healthy weight and potentially lower their medical expenses.
3. **Targeted Insurance Plans:** Design insurance plans that cater to different age groups. Tailored plans can address the unique healthcare needs of people at various life stages.
4. **Model Refinement:** Continue refining and optimizing predictive models. This could involve exploring more advanced algorithms or techniques to improve the accuracy of predictions.
5. **Regular Data Updates:** Regularly update the dataset with new data to ensure that the models remain relevant and accurate. Healthcare trends and cost patterns may change over time.
6. **Personalized Recommendations:** Provide personalized recommendations to individuals based on their age, smoking status, and BMI to help them make informed decisions about their healthcare and insurance choices.
7. **Collaboration with Healthcare Providers:** Collaborate with healthcare providers to gather more comprehensive data that includes detailed medical history. This could further enhance the accuracy of predictive models.
8. **Continuous Monitoring:** Keep monitoring the performance of predictive models and update them as necessary. New insights or changes in the healthcare landscape could require adjustments to the models.
9. **Customer Education:** Educate customers about how their lifestyle choices and age can impact their insurance costs. This transparency can help them make healthier decisions and potentially reduce their expenses.

By implementing these insights and suggestions, the insurance industry can make informed decisions to optimize insurance plans and promote healthier lifestyles among their customers.

Conclusion

This project demonstrates the process of predicting medical health insurance costs using various regression models. The models' performances are compared using evaluation metrics, and key insights are drawn from the analysis. The project highlights the importance of data preprocessing, model selection, and evaluation for accurate predictions in the field of medical insurance cost estimation. Further improvements and refinements could be explored to enhance the models' predictive capabilities.

Feel free to reach out for further discussions or clarifications! 😊

Contact Information

For any questions or inquiries, please contact:

Muhammad Habib

Email: muhammad.habib7@gmail.com

LinkedIn: [linkedin.com/in/muhammad-habib](https://www.linkedin.com/in/muhammad-habib)